Mindtree

Performance Test Report

For

Execution of

Kernel UIN Generator API – 100 users

Date: 28 April 2020

Author: Gaurav Sharan

Summary

This report presents the observations and findings of the load test conducted for a load of 100 users on kernel UIN generator service.

The objective of this load test was to observe and record behavior of the application when user load is increased from 20 to 100 in steps of 20.

Below are the scenario details:

Script/Report Name	Kernel UIN Generator
Run Date	27-April-2020
Period	10:53 UTC to 12:08 UTC
Number of concurrent users	20 to 100
Ramp up	01 users per second
Run Duration	----
Ramp down	NA

The transaction response times observed were as below:

Label	# Samples	Average(ms)	90% Line(ms)	Min(ms)	Max(ms)	Error %	Throughput
TR_kernel_generate-uin	5340	51467	81579	988	112198	0.00%	1.2/Sec

Performance Test Execution Details

We have executed JMeter script for kernel UIN generator service, which has transactions mentioned in the above table.

Average response time of the APIs is 5350 ms with 0% error rate.

Response time of the API is approx. 18 seconds when 20 users are active, it increases with increase in number of users. It reaches ~50 seconds when 60 users are active and gets to 82 second when 100 users are accessing the API.

TPS of the API is approx. 17.5 when 40 users are active and stays almost at same level. Sometimes it reaches 20.

Response Time and TPS Graph:

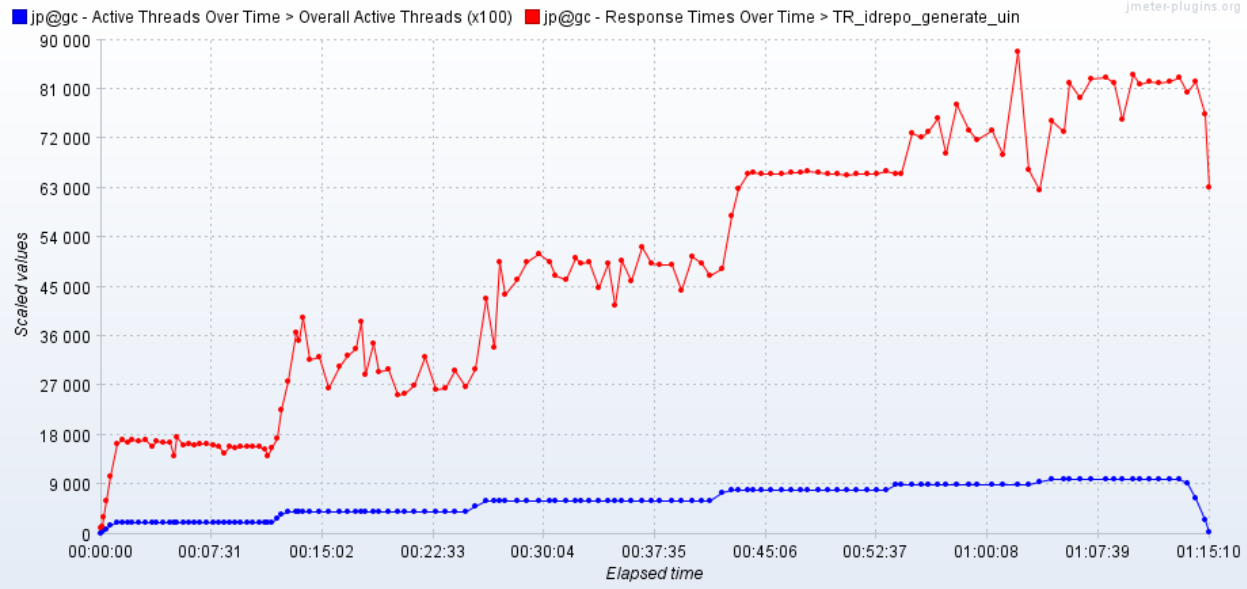

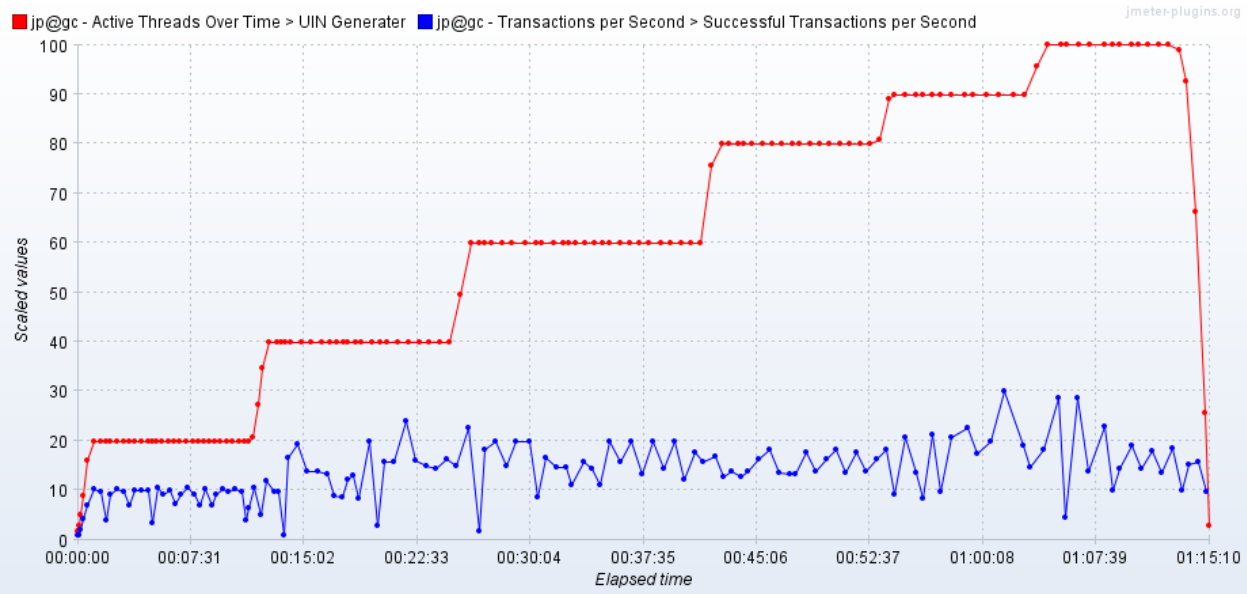

Mindtree

Resource Usage Pattern:

kernel Service cluster VM resource usage:

Average system load – 0.484 (4 cores)

Average CPU usage – 5 %

Memory used – 5 GB / 15.64 GB

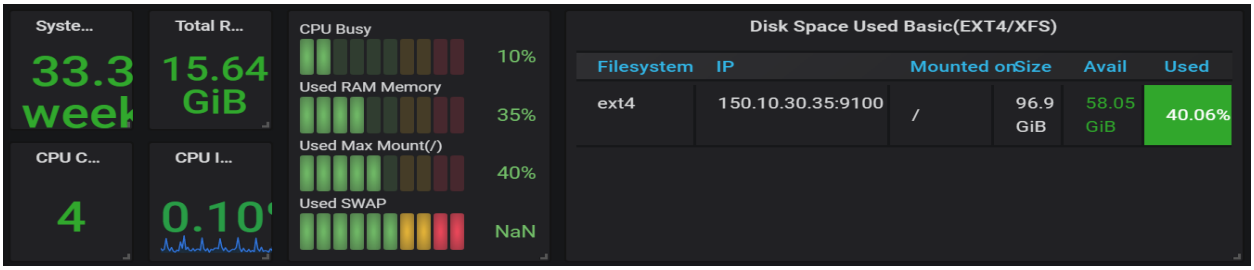

Kernel DB resource usage:

Average system load – 4.0 (2 cores) reaches max when 100 users arriveor say increases with increase in number of users

Average CPU usage – 91 %

Memory used – 1.4 GB / 7.78 GB